

Validity and Robustness of LLM-Assisted Policy Simulation in Health Services Research: A Value-Sensitivity Audit

Yichao Jin *

EPPS, University of Texas at Dallas

September 5, 2026

Abstract

Objective: To evaluate whether LLM-assisted policy simulations are valid, robust, and normatively transparent—or whether their outputs are sensitive to model choice, prompt framing, and latent value orientations. We audit the validity of LLM-generated policy simulations in health services research using administrative burden in U.S. vaccination access as a controlled testbed.

Study Setting and Design: We develop a multi-level diagnostic framework for measuring normative-frame sensitivity in LLM-generated policy simulations. Fifteen LLM endpoints spanning nine provider ecosystems complete a policy-value orientation battery and simulate vaccination willingness for 27 standardized synthetic profiles (controlled test cases) under low- and high-burden conditions in a stylized U.S. vaccination-access testbed. We estimate composite orientation effects (H1), narrative frame associations (H2), and within-model frame manipulation effects (H3).

Data Sources and Analytic Sample: Synthetic profiles (N=27) are experimental stimuli designed to isolate how models interpret administrative-burden scenarios, not to represent a human population. H1: 15 models; H2: 2,023 profile-repetition observations; H3: 108 model-profile clusters across 2,160 observations.

Principal Findings: (1) Composite orientation-effect: PVOC coefficient $\beta_1 = 1.51$, 95% CI $[-2.41, 5.43]$, $p = 0.45$ —composite orientation does not predict cross-model variation. (2) Narrative frame association: Access-barrier rationales associated with larger simulated effects ($\beta = 2.34$, 95% CI $[0.87, 3.81]$); collective-responsibility rationales associated with smaller effects ($\beta = -1.89$, 95% CI $[-3.12, -0.66]$). (3) Within-model manipulation: Autonomy frames increase simulated burden effects by 1.8 percentage points (95% CI $[0.4, 3.2]$); collective-obligation frames reduce effects by 2.1 percentage points (95% CI $[-3.5, -0.7]$).

Conclusions: LLM-assisted policy simulation is normatively frame-sensitive. The same model, same synthetic profiles, and same policy scenario can produce materially different simulated treatment-effect estimates—shifts of approximately 3–11 percentage points on a 0–100 scale—depending on the analyst-selected normative framing. H3 does not identify latent model values; it quantifies the sensitivity of simulated policy estimates to analyst-selected normative framing. This frame sensitivity is itself an operational vulnerability that validity audits must detect. Researchers

*University of Texas at Dallas, Yichao.Jin@UTDallas.edu

should audit simulation validity before deploying LLMs as policy-simulation tools in HSR.

What is known on this topic

- LLMs are increasingly used to simulate social judgments and policy effects in health services research.
- Prior work measures model values through direct survey prompts rather than downstream simulated policy effects.
- Less is known about whether latent value orientations shape simulated treatment effects under policy conditions.

What this study found

- Composite policy-value scores do not predict cross-model variation in simulated burden effects.
- Narrative interpretive frames (access-barrier, collective-responsibility) are systematically associated with simulated effect sizes.
- Within-model frame manipulation confirms: autonomy and equity/access frames increase simulated burden effects; collective-obligation frames reduce them.
- LLM-assisted policy simulation requires systematic validity auditing before deployment in HSR.

Keywords: large language models, simulation validity, administrative burden, policy simulation robustness, normative-frame sensitivity, calibration

1 Introduction

Administrative burden in U.S. health services—the learning, psychological, and compliance costs that individuals face when interacting with public programs—is a well-documented barrier to care [7]. These burdens include forms, documentation requirements, waiting time, procedural uncertainty, and perceived loss of autonomy, and they can delay or discourage access to vaccination and other preventive services [7]. Understanding how administrative burden shapes health-service access is a central question in U.S. health services research (HSR), where policymakers need tools to anticipate the likely effects of procedural changes before implementation.

Large language models (LLMs) are increasingly used in HSR not only to generate text but also to simulate social judgments, preferences, and behavioral responses [1–4]. When applied to health-service delivery, an LLM may be asked to simulate how patients or program beneficiaries would respond to a policy intervention. This creates a methodological chain:

administrative burden → health-service access/delivery → LLM policy simulation → validity threat from

The core validity question is whether LLM-assisted policy simulations are robust across models and framing choices. Vaccination access serves as a controlled HSR testbed for evaluating this validity threat. This paper evaluates whether simulated treatment effects are sensitive to the normative lens activated by model architecture, training data, or prompt design.

A central challenge is that model provenance is an inadequate proxy for value orientation. To avoid treating development origin as a fixed ideological category, I independently

measure each model’s policy-value orientation using a structured battery capturing five dimensions: autonomy protection, state trust, collective obligation, administrative-burden sensitivity, and coercion aversion.

The study addresses three research questions. First, do LLMs with different empirically measured policy-value orientations produce different simulated marginal effects of administrative burden on vaccination willingness? Second, are simulated burden effects systematically associated with the interpretive frames used in model rationales? Third, can simulated burden effects be shifted within the same base model by experimentally assigning governance-value frames?

Contributions: (1) shift LLM value-sensitivity study from opinion to simulated policy effects; (2) introduce a policy-value orientation battery; (3) provide a multi-level diagnostic framework for LLM simulation validity.

2 Theoretical Framework and Hypotheses

2.1 Policy-Value Orientation

Policy-value orientation is defined as systematic tendencies to interpret governance situations through assumptions about individual–institution relationships, legitimacy of public authority, and procedural costs. The framework captures five dimensions (autonomy protection, state trust, collective obligation, burden sensitivity, coercion aversion) via a structured battery. The PVOC aggregates these dimensions with theoretically specified weights, treated as a diagnostic aggregate rather than validated scale. Internal consistency is mixed (only collective obligation reaches $\alpha = 0.869$), consistent with prior findings that human psychometric scales may not transfer directly to LLMs.

2.2 Hypotheses

The empirical tests proceed in three linked steps, moving from an aggregate orientation index to narrative-level associations to within-model causal evidence. The tests form a diagnostic framework rather than a single confirmatory chain: each step is informative regardless of the outcome of the preceding one.

H1 (Composite Orientation–Effect). If LLM policy simulations are systematically influenced by latent policy-value orientations, then models with higher composite PVOC scores (Policy-Value Orientation Composite, indexing autonomy protection, burden sensitivity, and coercion aversion against state trust and collective obligation) should produce larger simulated administrative-burden effects on vaccination willingness. A null result would imply that value sensitivity, if present, does not follow a simple composite pathway, suggesting the need for more granular diagnostics such as narrative frame decomposition or experimental manipulation.

H2 (Narrative Frame Association). Interpretive frames expressed in model rationales are expected to be systematically associated with simulated burden effects, independent of the composite PVOC. Access-barrier, autonomy-infringement, and coercive-backlash frames are expected to correspond to larger simulated burden effects; procedural-legitimacy and collective-responsibility frames are expected to correspond to smaller effects. Because frames are generated by the model rather than experimentally assigned, H2 identifies descriptive associations rather than causal pathways.

H3 (Frame Manipulation). Randomly assigning governance-value frames to the same base model is expected to causally shift simulated burden effects relative to a neutral frame. Autonomy and equity/access frames are expected to increase simulated burden effects, while collective-obligation frames are expected to decrease them. Because the base model, profiles, scenarios, and output format are held fixed, any shift in simulated effects can be attributed to the frame itself rather than to stable model characteristics.

The three hypotheses are designed as a diagnostic framework rather than a sequential gamble. Each addresses a distinct question—aggregate orientation, narrative association, and causal frame—and each is informative whether or not the others yield significant results.




Figure 1: Diagnostic framework for three levels of value sensitivity testing. H1 tests whether a model-level PVOC predicts cross-model variation in administrative-burden effects. H2 tests whether narrative frames expressed in model rationales are systematically associated with predicted effects. H3 tests whether experimentally assigned governance-value frames causally shift predicted effects within the same base model. The three tests are designed as independent diagnostic components: each is informative regardless of the outcome of the others.

3 Experimental Design and Methods

3.1 Model Sample

We evaluate fifteen LLM endpoints spanning nine provider ecosystems (OpenAI, Anthropic, Google, Meta, Mistral, Cohere, Moonshot AI, DeepSeek, Qwen). The sample includes frontier proprietary, lower-cost, hosted open-weight, enterprise-oriented, and non-U.S. models. All models were run with temperature = 0 where supported; each cell repeated five times. Detailed identifiers are reported in the Online Supplement.

3.2 Policy-Value Orientation Battery

We administered a 20-item Likert battery measuring five dimensions: autonomy protection, state trust, collective obligation, administrative-burden sensitivity, and coercion aversion. Each dimension uses four items on a 1–7 scale. The Policy-Value Orientation Composite (PVOC) aggregates the five dimensions with theoretically specified weights, treated as a diagnostic aggregate rather than validated scale.

3.3 Administrative-Burden Simulation

To measure each model’s simulated burden effect, we designed a vaccination-access simulation as a controlled HSR testbed, in which the same standardized profile faces two policy conditions: low administrative burden and high administrative burden. The simulation was administered after the orientation battery but before any frame manipulation, using a fresh API call for each model–profile–condition–repetition.

3.3.1 Standardized synthetic profiles (controlled test cases)

We constructed 27 standardized synthetic profiles as controlled test cases that span theoretically relevant variation in income, work flexibility, and health-care access while holding age, gender, residence, education, family status, and trust in public health authorities constant. These profiles are not intended to represent a human population; they are experimental stimuli designed to isolate how models interpret the same administrative-burden scenarios under different model and frame conditions. The full profile set is reported in the Online Supplement.

3.3.2 Low- and high-burden scenarios

Two burden conditions differ in administrative requirements: low-burden (online scheduling, nearby pharmacy, minimal documentation, 15 minutes) and high-burden (in-person registration, income verification, proof of residency, paper forms, 1–2 hour wait, one-week advance appointment). These scenarios serve as a **stylized testbed** for auditing LLM simulation validity, not as estimates of real-world burden magnitude. Full prompts are in the Online Supplement.

3.3.3 Outcome: predicted willingness and marginal burden effect

For each model, profile, burden level, and repetition, we obtained a predicted vaccination willingness on a 0–100 scale. The model produced a structured JSON output containing an integer willingness score and a brief rationale. The primary outcome is the marginal

effect of moving from low to high burden for a given model, profile, and repetition: $\Delta_{mir} = W_{mi,low,r} - W_{mi,high,r}$. Positive values indicate that the model predicts a reduction in willingness under higher burden. For model-level analyses, we average across profiles and repetitions.

3.4 Task-Compliance Diagnostics

Because comparable general benchmark scores are not uniformly available for all fifteen model endpoints, we do not construct a composite capability score or include capability-adjusted regressions. Instead, we assess whether each model complied with the simulation protocol using task-specific diagnostics computed from the study’s own outputs: JSON validity, score-range compliance, non-refusal rate, repetition-level stability (within-profile-burden standard deviation across five repetitions), burden-effect dispersion (standard deviation of profile-level effects), and flip rate. These diagnostics, reported in Table 3, verify protocol compliance and output stability rather than construct a general capability ranking. The substantive analyses therefore focus on measured policy-value orientation, predicted burden effects, narrative frames, and within-model frame responsiveness.

3.5 Narrative coding framework

To test H2, we coded rationales using five theoretically derived frames: autonomy infringement, procedural legitimacy, collective responsibility, access barriers, and coercive backlash. A rule-based algorithm with predefined keyword patterns (regular expressions) coded each rationale; frames are non-mutually exclusive. Validation on 200 rationales (10%) confirmed pattern validity. Binary frame indicators are used as descriptive evidence of rationale–outcome co-expression. Full patterns are in the Online Supplement.

3.6 Within-Model Frame Manipulation

To test H3, we conducted a within-model experiment varying system-prompt frames while holding base model, profiles, scenarios, and output format fixed. This provides within-model evidence that framing influences simulated burden effects.

3.6.1 Model selection

The frame-manipulation experiment was conducted on a purposive four-model subset (GPT-4.1, Llama 3.1 70B Instruct, Qwen3.6-72B Instruct, and Mistral Large) representing a frontier proprietary closed-source model, a hosted open-weight model, a non-U.S. model ecosystem, and a low-PVOC robustness case. This design tests whether frame sensitivity can be induced within diverse model endpoints, not the distribution of frame responsiveness across all fifteen models.

3.6.2 Frame conditions and randomization

Four system-prompt frames were used: *neutral*, *autonomy* (weighting individual choice, privacy, and procedural dignity), *collective obligation* (weighting civic responsibility and public health coordination), and *equity/access* (weighting unequal procedural costs and

barriers faced by disadvantaged groups). The manipulation was fully crossed: each model evaluated every profile under each frame, with frame order randomized across runs.

3.6.3 Statistical analysis

To test whether the assigned frame shifts Δ , we estimate the burden effect with model and profile fixed effects, and indicator variables for each value frame (neutral as the reference category), adjusting pairwise comparisons using the Holm–Bonferroni correction. Because the frame manipulation is fully crossed within each model, profile fixed effects are identified within-model and absorb all time-invariant profile-level heterogeneity. Standard errors are clustered at the profile level to account for the five repeated queries per profile, treating the profile as the primary sampling unit. The inferential target is the profile-level frame contrast, not a population-level treatment effect. Full results are reported in Section 4.4.3 and the Online Supplement.

3.7 Estimation Strategy

We test H1–H3 using hypothesis-specific estimation strategies. Because the main explanatory variables in H1 and H2 vary only at the model level, these analyses are interpreted cautiously: the primary emphasis is on coefficient direction, magnitude, confidence intervals, and robustness across alternative specifications rather than on conventional significance thresholds alone.

3.7.1 Orientation-effect models (H1)

The primary H1 analysis tests whether models with higher PVOC scores predict larger burden effects:

$$\bar{\Delta}_m = \beta_0 + \beta_1 \text{PVOC}_m + \epsilon_m,$$

where $\bar{\Delta}_m$ is the average burden effect across profiles and repetitions. Because the identifying variation in PVOC exists only at the model level, this analysis is interpreted as a descriptive cross-model association rather than a high-powered population estimate. We also estimate separate bivariate regressions replacing the composite PVOC with each of the five orientation dimensions (autonomy protection, state trust, collective obligation, administrative-burden sensitivity, and coercion aversion) as diagnostic checks, not as confirmatory tests of validated subscales.

3.7.2 Narrative frame-decomposition models (H2)

The H2 analysis decomposes associations between PVOC, expressed rationale frames, and predicted burden effects at the model–profile level. We refer to this as “mediation-style” because PVOC is measured rather than randomized and varies only at the model level; the estimates should be interpreted as frame–outcome association diagnostics rather than formal causal mediation. The H2 product terms summarize whether models with different PVOC scores tend to express different frames and whether those expressed frames are associated with different predicted burden effects. H2 is evaluated primarily by the direction and stability of the frame–outcome associations; because each model–profile combination contributes five repeated runs, standard errors are clustered at the model–profile level, and frame-specific p -values are adjusted using the Benjamini–Hochberg false discovery rate procedure at $q = 0.10$.

3.7.3 Frame-manipulation models (H3)

Using the within-model frame-manipulation data, we estimate a two-way fixed-effects model:

$$\Delta_{mcir} = \alpha_m + \mu_i + \beta_1 \mathbb{I}(c = \text{Autonomy}) + \beta_2 \mathbb{I}(c = \text{Collective}) + \beta_3 \mathbb{I}(c = \text{Equity}) + \varepsilon_{mcir},$$

with neutral as the omitted reference, model fixed effects α_m , and profile fixed effects μ_i . Pairwise comparisons against the neutral frame are adjusted using the Holm–Bonferroni correction, and model-specific frame effects are reported separately. Because each model–profile–frame cell contains five repeated runs, standard errors are clustered at the model–profile level to account for within-cell correlation across repetitions.

3.7.4 Multiple-testing and robustness checks

Within hypothesis families, we apply the Benjamini–Hochberg FDR procedure at $q = 0.10$ for H2 (five frames) and the Holm–Bonferroni correction for H3 (three value frames). H1’s primary test concerns a single coefficient, so no adjustment is applied. Robustness checks exclude each synthetic profile one at a time; replace PVOC with its five separate dimension scores; use an unweighted raw-score PVOC; use binary frame presence rather than intensity in H2; and exclude models below alternative willingness thresholds.

4 Results

4.1 Data Quality and Descriptive Statistics

4.2 Predicted Administrative-Burden Effects

4.3 Orientation Profiles and Burden Effects

Table 1 reports policy-value orientation profiles and predicted administrative-burden effects by model. Higher PVOC values indicate profiles expected to generate larger burden effects. For every model, mean predicted willingness is lower under high-burden conditions. Variation is not reducible to baseline willingness, motivating analyses of whether cross-model variation is associated with policy-value orientation and rationale frames.

4.4 Hypothesis Tests

This subsection reports the hypothesis-specific tests as three independent diagnostic components consistent with the framework introduced earlier. H1 examines whether the model-level PVOC predicts cross-model variation in average administrative-burden effects. H2 examines whether narrative frames expressed in model rationales are systematically associated with predicted burden effects, independently of any aggregate orientation score. H3 examines whether experimentally assigned governance-value frames shift predicted burden effects within the same base model. The three tests are independently informative: each addresses a distinct level of evidence (aggregate orientation, narrative mechanism, and within-model causal frame), and each can contribute to the overall assessment whether or not the others yield significant results.

Table 1: Policy-value orientation profiles and predicted administrative-burden effects by model

Model	PVC	Low W	High W	$\bar{\Delta}$	95% CI
Llama 3.1 8B Instruct	2.90	84.33	46.48	37.85	[37.06, 38.65]
Claude Haiku	2.20	70.96	42.74	28.22	[27.13, 29.32]
Llama 3.1 70B Instruct	1.85	81.11	48.30	32.81	[31.16, 34.47]
Claude Sonnet 4.6	1.21	71.13	32.09	39.04	[38.06, 40.01]
Qwen3.6-72B Instruct	0.75	78.96	44.19	34.78	[33.29, 36.27]
Gemini 2.5 Flash	0.31	88.01	29.81	58.19	[56.57, 59.81]
Gemini 2.5 Pro	0.01	87.93	23.07	64.85	[63.69, 66.01]
GPT-4.1 mini	-0.24	76.48	41.44	35.04	[34.28, 35.80]
Cohere Command R+	-0.28	75.11	45.63	29.48	[28.49, 30.46]
Claude Opus 4.8	-0.41	70.73	35.33	35.40	[34.35, 36.45]
GPT-4.1	-0.61	82.81	53.78	29.04	[27.82, 30.26]
Kimi K2.6	-0.84	77.07	39.57	37.51	[36.41, 38.62]
DeepSeek-V3.2	-1.30	76.63	27.52	49.11	[48.23, 50.00]
GPT-4o	-1.89	83.56	59.67	23.89	[22.68, 25.10]
Mistral Large	-3.65	71.16	50.30	20.87	[19.92, 21.82]

Notes: PVC = Policy-Value Orientation Composite, standardized across the fifteen-model sample. Low W and High W are mean predicted vaccination willingness scores under low- and high-burden conditions. $\bar{\Delta}$ = Low W - High W . 95% CI calculated as $\bar{\Delta} \pm 1.96 \times SE$. Table sorted by PVC in descending order.

4.4.1 H1: Orientation–Effect Models

H1 tests whether PVOC predicts simulated burden effects using a bivariate model-level regression ($N=15$). The estimated coefficient is positive but small ($\hat{\beta} = 1.51$, 95% CI $[-2.41, 5.43]$, $R^2 = 0.044$), not reaching conventional detectability thresholds (minimum detectable $R^2 \approx 0.40$). The composite PVOC does not provide evidence of stable linear association with predicted burden effects.

Figure 2 visualizes the bivariate relationship. Consistent with the regression result, the fitted line is shallow and the points are widely dispersed. Models with higher PVOC scores do not necessarily predict larger burden effects: Llama 3.1 8B Instruct has the highest PVOC score but predicts a mid-range burden effect, while Gemini 2.5 Pro has a near-zero PVOC score but predicts the largest effect in the sample.

The null H1 result should be interpreted in light of three limitations. First, the PVOC imposes strong theoretical assumptions—equal weighting across five dimensions, fixed directional signs, and an additive structure—that may not match how value dimensions interact with simulated policy effects. Second, as documented in Section ??, the PVOC has limited internal consistency across its constituent dimensions, with only collective obligation reaching conventional reliability thresholds ($\alpha = 0.869$). A composite aggregate whose components do not cohere as a latent construct cannot be expected to produce stable cross-model predictions. Third, at $N = 15$ model endpoints, model-level regressions have limited statistical power, making a null composite result difficult to interpret in isolation [5, 6]. The null composite result therefore rules out a simple additive pathway but does not constitute evidence against normative-frame sensitivity through other, more granular channels. This limitation reinforces the value of the dimension-specific, narrative-level, and experimental evidence that follows.

This null result is substantively informative. Models vary considerably in both measured policy-value orientation and predicted burden effects, but the composite PVOC does not explain that cross-model variation. This negative finding is valuable: a scalar ideology/value score is insufficient to audit LLM simulation; frame-level diagnostics are more sensitive. The expanded sample strengthens the conclusion that simulated burden effects are not well captured by a single aggregate policy-value score; subsequent analyses examine whether more specific expressed frames and experimentally assigned governance-value lenses provide stronger evidence of normative-frame sensitivity in LLM-assisted policy simulation.




Figure 2: Policy-value contamination index and predicted administrative-burden effects

Notes: Each point represents one model in the expanded fifteen-model sample. The horizontal axis reports the model's PVOC, and the vertical axis reports the model-level average burden effect, $\bar{\Delta}_m = \bar{W}_m^{low} - \bar{W}_m^{high}$. The dashed line shows the fitted bivariate regression. The shallow fitted line and low $R^2 = 0.044$ are consistent with the null H1 result.

4.4.2 H2: Narrative Frame-Decomposition Analysis

H2 examines whether interpretive frames in model rationales are associated with predicted administrative-burden effects, independently of any aggregate orientation index. The analysis is designed as a frame–outcome association decomposition: the goal is not to establish a causal pathway from PVC to burden predictions, but to assess whether expressed rationale frames are systematically related to predicted burden effects in theoretically interpretable ways.

We first examined the prevalence of the five theoretically derived frames in the high-burden rationales. Three frames displayed meaningful variation and were retained in the expanded fifteen-model frame-decomposition analysis. Access-barrier reasoning appeared in 52.7% of high-burden rationales after recoding the frame to require an explicit link between procedural costs and reduced access, resource constraints, work flexibility, child-care, provider access, or socioeconomic vulnerability. Collective-responsibility reasoning appeared in 59.2% of rationales, while coercive-backlash reasoning appeared in 8.4%. Autonomy infringement and procedural legitimacy showed little or no meaningful variation and were therefore reported descriptively but excluded from the main frame model.

Table 2 reports the retained frame–outcome estimates. Access-barrier reasoning is associated with larger predicted burden effects ($\hat{b} = 8.83$, $p < 0.001$), indicating that rationales linking administrative requirements to concrete access constraints correspond to larger predicted reductions in willingness. Collective-responsibility reasoning is associated with smaller burden effects ($\hat{b} = -6.20$, $p < 0.001$), suggesting that rationales emphasizing civic duty or public health responsibility dampen the predicted deterrent effect of administrative burden. Coercive-backlash reasoning is also associated with larger burden effects ($\hat{b} = 6.02$, $p < 0.001$), consistent with the idea that resistance or distrust frames amplify predicted burden sensitivity. Because each model–profile combination contributes five repeated runs, standard errors are clustered at the model–profile level to account for within-cell correlation.

Leave-one-model-out diagnostics support the stability of the frame–outcome associations: access-barrier reasoning remains positive in all fifteen leave-one-out samples, collective-responsibility reasoning remains negative in all fifteen, and coercive-backlash reasoning remains positive in fourteen of fifteen. The direct PVOC coefficient in the frame-adjusted model is positive and sign-stable ($\hat{b} = 1.55$, $p < 0.001$) but interpreted cautiously because PVOC is measured rather than randomized and varies only at the model level. Overall, H2 receives partial support: the evidence shows that model rationales contain interpretable frames that move systematically with simulated burden effects, consistent with value sensitivity through expressed narrative frames rather than a single aggregate orientation score.

4.4.3 H3: Within-Model Frame Manipulation

H3 tests whether governance-value frames can shift predicted burden effects within the same base model, using a purposive four-model subset: GPT-4.1, Llama 3.1 70B Instruct, Qwen3.6-72B Instruct, and Mistral Large.

Figure 3 reports pooled frame effects, using the neutral frame as the reference category. The inferential unit is the profile–frame combination, with standard errors clustered at the profile level to account for the five repeated queries per profile. The results strongly support H3: the autonomy frame increases the predicted burden effect by 2.62 points (95% CI [1.95, 3.29]); the collective-obligation frame decreases it by 10.54 points (95% CI [-

Table 2: Narrative frame-decomposition estimates

Frame	Prevalence	b : Frame $\rightarrow \Delta$	SE	p	LOMO sign stable
Access barriers	52.7%	8.83	0.54	< 0.001	Yes
Collective responsibility	59.2%	-6.20	0.54	< 0.001	Yes
Coercive backlash	8.4%	6.02	0.97	< 0.001	Yes

Notes: The table reports descriptive frame–outcome estimates using high-burden rationales in the expanded fifteen-model sample. The dependent variable is the profile–repetition-level burden effect $\Delta = W^{low} - W^{high}$. Coefficients are from a frame-adjusted model including PVC and the three retained frame indicators. $N = 2,023$, $R^2 = 0.234$. Positive coefficients indicate larger predicted burden effects. LOMO = leave-one-model-out. Access barriers and collective responsibility are sign-stable in all fifteen leave-one-model-out samples; coercive backlash is positive in fourteen of fifteen. The table emphasizes frame–outcome associations rather than formal mediation. PVC-to-frame paths and $a \times b$ products are not interpreted as causal mediation estimates because PVC varies only at the model level and is not experimentally assigned.

11.22, -9.87]); the equity/access frame increases it by 6.29 points (95% CI [5.61, 6.96]). All three frame contrasts are statistically significant at conventional levels. The pooled four-model specification explains substantial variation in burden effects ($N = 2,160$ profile–repetition observations, 108 unique profiles, adjusted $R^2 = 0.66$).

These estimates should be interpreted as profile-level frame contrasts rather than population-level treatment effects. The 2,160 observations are not independent patient-level observations; they represent 27 synthetic profiles $\times$ 4 frames $\times$ 4 models, with five repetitions per profile–frame–model cell. The clustering adjustment treats the profile as the primary sampling unit, reflecting the study design in which synthetic profiles are the controlled test cases and model queries are repeated measurements. The small standard errors (SE ≈ 0.34) and narrow confidence intervals reflect the within-profile consistency of LLM responses at temperature = 0, not population-level precision in the traditional sense. The effect sizes (2.62, 6.29, 10.54 points on a 0–100 willingness scale) are the primary quantities of interest, with uncertainty quantified via profile-clustered standard errors.

These effects are directionally consistent with the theoretical expectations: the autonomy frame increases burden sensitivity by emphasizing individual choice, privacy, and procedural dignity; the collective-obligation frame reduces it by emphasizing civic responsibility and tolerance of inconvenience; the equity/access frame increases it by emphasizing unequal procedural costs. The Mistral Large extension is especially informative: although Mistral Large has the lowest PVOC score and the smallest baseline burden effect, all three frame manipulations shift its simulated burden effects in the expected directions. Taken together, H3 provides the clearest evidence that governance-value framing can shape LLM-assisted policy simulation: although H1 did not find a stable cross-model composite association, the within-model manipulation shows that activating specific governance-value frames is sufficient to shift simulated treatment effects in theoretically predictable directions.

4.5 Robustness Checks

For H1, the expanded fifteen-model analysis reinforces the null result for the PVOC ($\hat{\beta} = 1.51$, 95% CI $[-2.41, 5.43]$, $R^2 = 0.044$). Dimension-specific checks also caution

Figure 3: Forest plot of within-model frame-manipulation effects

		Frame contrast	Coefficient [95% CI]		
Pooled estimates (four-model average):		Collective obligation vs. neutral	−10.54 [−11.22, −9.87]		
		Equity/access vs. neutral	6.29 [5.61, 6.96]		
		Autonomy vs. neutral	2.62 [1.95, 3.29]		
		Model	Autonomy	Collective	Equity
Model-specific estimates:		GPT-4.1	2.41 [1.52, 3.30]	−9.78 [−10.67, −8.89]	5.93 [5.04, 6.82]
		Llama 3.1 70B	2.96 [2.07, 3.85]	−11.19 [−12.08, −10.30]	5.52 [4.63, 6.41]
		Qwen3.6-72B	2.48 [1.59, 3.37]	−10.37 [−11.26, −9.48]	7.41 [6.52, 8.30]
		Mistral Large	2.63 [1.74, 3.52]	−10.81 [−11.70, −9.92]	6.30 [5.41, 7.19]

Notes: Coefficients represent the shift in predicted burden effect (percentage points on 0–100 willingness scale) relative to the neutral frame. Positive values indicate larger predicted burden effects. Pooled estimates use model and profile fixed effects with standard errors clustered at the profile level ($N = 2,160$, adjusted $R^2 = 0.66$). Model-specific estimates show consistent directional patterns across all four base models.

against overinterpreting the original nine-model results: the initial collective-obligation signal weakens and is no longer statistically significant after sample expansion, and is best interpreted as a sample-sensitive exploratory pattern rather than a stable dimension-level predictor. For H2, leave-one-model-out diagnostics show stable frame coefficients: access-barrier reasoning remains positive in all fifteen leave-one-out samples, collective-responsibility reasoning remains negative in all fifteen, and coercive-backlash reasoning remains positive in fourteen of fifteen. For H3, model-specific frame regressions show the pooled manipulation is not driven by a single base model: the collective-obligation frame reduces predicted burden effects in all four models, and the autonomy frame is positive in all four. The equity/access frame is more heterogeneous (positive for GPT-4.1 and Qwen3.6-72B, no effect for Llama 3.1 70B). Detailed H1, H2, and H3 robustness results are reported in the Online Supplement.

Robustness checks support granular interpretation: simulated policy effects are more sensitive to interpretive frames and assigned governance-value frames than to a single aggregate PVOC score.

5 Discussion

This study examined LLM policy simulation validity using administrative burden in vaccination access as a controlled HSR testbed. Results support frame-sensitive validity: the aggregate PVC index does not predict simulated burden effects across models, suggesting value sensitivity is not captured by a single scalar index.

The H1 null result is robust to alternative constructions, indicating that value sensitivity is not well captured by a single scalar index.

The narrative analysis provides more stable rationale–outcome co-expression evidence. Access-barrier and coercive-backlash rationales are associated with larger predicted burden effects; collective-responsibility rationales with smaller effects. These associations remain stable in leave-one-model-out diagnostics. H2 should be understood as descriptive

co-expression rather than causal mediation: willingness scores and rationales are generated simultaneously, so observed associations may reflect post-hoc justification rather than a causal mechanism.

The strongest evidence comes from the within-model frame manipulation. In a purposive four-model subset, assigned governance-value frames shifted simulated burden effects in expected directions: autonomy and equity/access frames increased effects; collective-obligation reduced them. This shows that LLM-assisted policy simulation is normatively frame-sensitive: the same model, profiles, and scenarios can produce materially different estimates depending on the normative lens activated by the prompt.

The H3 frame-manipulation results may reflect instruction-following, but this is precisely the operational mechanism of frame sensitivity. H3 quantifies steerability under alternative governance-value frames—a swing of 3–11 points on a 0–100 scale—which is the fragility that simulation-validity audits are designed to detect.

The findings have three implications. First, audits should move beyond stated-opinion tests to examine how models simulate policy effects under different frames. Second, audits should measure specific dimensions and frame sensitivity rather than rely on coarse provenance labels. Third, prompt framing should be treated as part of the experimental design: simulations should not be reported as single neutral estimates when governance-value frames materially shift predicted effects. Analysts should report model identity, prompt framing, scenario wording, and robustness to alternative frames.

5.1 Limitations and Future Research

Six limitations: (1) small sample ($N=15$) limits H1 power; (2) H1–H3 are independent diagnostics; (3) bundled burden manipulation; (4) synthetic profiles are controlled cases; (5) coding requires validation; (6) observations are profile-clustered, with small SEs reflecting LLM consistency at temperature=0.

Seventh, and most importantly for health services research, the vaccination-access scenario is a **stylized testbed** rather than a realistic representation of any specific U.S. vaccination program. The income verification and residency requirements in the high-burden condition are stylized procedural barriers typical of means-tested public programs, not realistic requirements for vaccination access. The scenario was designed to create a controlled contrast for testing LLM simulation validity, not to estimate the magnitude of administrative burden in actual health service delivery. The findings should therefore be interpreted as evidence about LLM-assisted policy simulation methodology—specifically, its sensitivity to normative framing—rather than as estimates of real-world vaccination access barriers. Future work should validate the diagnostic approach in realistic HSR settings with externally calibrated human-subject data. validity rather than as evidence about real-world administrative burden effects; vaccination access serves as a controlled HSR testbed, and the validity and robustness conclusions generalize to LLM-assisted policy simulation methodology, not to the specific behavioral dynamics of vaccination uptake. Future work could integrate the diagnostic approach with human-subject data to calibrate whether value sensitivity in LLM simulations mirrors or diverges from value sensitivit... [truncated]

5.2 Conclusion

The central finding is normative-frame sensitivity: expressed narrative frames are systematically associated with simulated burden effects, and experimentally assigned governance-value frames shift simulated treatment effects within the same model. Validity depends less on aggregate orientation scores than on interpretive framing.

For researchers and policymakers, the implication is clear: LLM-assisted policy simulations should not be treated as neutral behavioral forecasts without auditing their sensitivity to model choice, prompt framing, rationale structure, and scenario design. The goal is not to eliminate values from policy simulation, but to make their influence measurable, transparent, and contestable—so that the method can be used responsibly as an HSR tool. This recommendation is consistent with broader concerns about the risks and institutional consequences of foundation-model deployment .

References

- [1] Gati Aher, Rosa I. Arriaga, and Adam Tauman Kalai. Using large language models to simulate multiple humans and replicate human subject studies. *arXiv preprint arXiv:2208.10264*, 2022. URL <https://arxiv.org/abs/2208.10264>.
- [2] Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, 2023. doi: 10.1017/pan.2023.2.
- [3] Zhicheng Lin. Large language models as linguistic simulators and cognitive models in human research. *arXiv preprint arXiv:2402.04470*, 2024. URL <https://arxiv.org/abs/2402.04470>.
- [4] Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*. Association for Computing Machinery, 2023. doi: 10.1145/3586183.3606763.
- [5] Tai-Quan Peng, Kaiqi Yang, Sanguk Lee, Hang Li, Yucheng Chu, and Yu-Ping Lin. Beyond partisan leaning: A comparative analysis of political bias in large language models. *Journal of Information Technology & Politics*, 23(2):145–168, 2026. doi: 10.1080/19331681.2026.2646990. Published online March 2026.
- [6] Paul Röttger, Valentin Hofmann, Valentina Pyatkin, Musashi Hinck, Hannah Rose Kirk, Hinrich Schütze, and Dirk Hovy. Political compass or spinning arrow? towards more meaningful evaluations for values and opinions in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15295–15311, Bangkok, Thailand, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.816. URL <https://aclanthology.org/2024.acl-long.816/>.
- [7] Jae Y. Shin et al. Patient administrative burden: A scoping review. *Health Affairs Scholar*, 3(11):qxaf216, 2025. doi: 10.1093/haschl/qxaf216.

Online Supplement: “Validity and Robustness of LLM-Assisted Policy Simulation in Health Services Research: A Value-Sensitivity Audit” Yichao Jin September 5, 2026

This Online Supplement contains supporting materials for the main text, organized as follows. The full text of each appendix is provided in the same order as in the original submission.

6 Data Quality and Simulation Diagnostics

Table 3 reports task-compliance diagnostics for the expanded fifteen-model simulation (4,050 raw responses, 2,025 low–high burden-effect pairs). JSON validity, score-range compliance, and non-refusal rates were 100% for thirteen of fifteen models; Cohere Command R+ and Kimi K2.6 each had one malformed response (99.6% compliance). The flip rate was 0% for every model, indicating that all models recognized the direction of the burden manipulation. Repetition-level within-profile–burden standard deviations ranged from 0.00 to 3.18 points on the 0–100 scale, indicating high within-model stability.

Table 3: Data quality and simulation diagnostics by model

Model	<i>N</i>	JSON%	Score%	NonRef%	Mean <i>W</i>	StabSD	DeltaSD	Flip%
GPT-4o	270	100.0	100.0	100.0	71.61	0.43	7.13	0.0
GPT-4.1	270	100.0	100.0	100.0	68.30	0.62	7.12	0.0
Llama 3.1 8B	270	100.0	100.0	100.0	65.41	0.28	4.54	0.0
Llama 3.1 70B	270	100.0	100.0	100.0	64.70	0.37	9.75	0.0
Qwen3.6-72B	270	100.0	100.0	100.0	61.57	3.18	7.22	0.0
Mistral Large	270	100.0	100.0	100.0	60.73	0.88	5.41	0.0
Cohere Command R+	270	99.6	99.6	100.0	60.32	1.52	4.77	0.0
GPT-4.1 mini	270	100.0	100.0	100.0	58.96	0.87	3.86	0.0
Gemini 2.5 Flash	270	100.0	100.0	100.0	58.91	1.57	9.05	0.0
Kimi K2.6	270	99.6	99.6	100.0	58.25	1.39	5.83	0.0
Claude Haiku	270	100.0	100.0	100.0	56.85	0.00	6.59	0.0
Gemini 2.5 Pro	270	100.0	100.0	100.0	55.50	2.49	5.83	0.0
Claude Opus 4.8	270	100.0	100.0	100.0	53.03	0.77	6.15	0.0
DeepSeek-V3.2	270	100.0	100.0	100.0	52.07	0.96	4.48	0.0
Claude Sonnet 4.6	270	100.0	100.0	100.0	51.61	0.46	5.72	0.0

Notes: *N* reports the number of raw simulation responses per model. JSON% is the percentage of outputs parsed as valid JSON. Score% is the percentage of responses with an integer willingness score in the 0–100 range. NonRef% is the percentage of responses that were not refusals or irrelevant answers. Mean *W* is the average predicted willingness across all profiles, burden conditions, and repetitions. StabSD is the mean within-profile–burden standard deviation across five repetitions. DeltaSD is the standard deviation of profile-level burden effects across the 27 synthetic profiles. Flip% is the percentage of profiles for which the average predicted willingness was higher under the high-burden condition than under the low-burden condition.

A Model Versioning, API Endpoints, and Decoding Parameters

This appendix reports the model identifiers, access routes, endpoints, and decoding settings used in the experiments. The study includes fifteen LLM endpoints. The original nine-model sample was accessed through official provider APIs or hosting endpoints where available; the six expansion models were accessed through OpenRouter using explicitly specified model identifiers. For OpenRouter-accessed models, results are interpreted as endpoint-level behavior because gateway access may involve routing, backend provider differences, or wrapper effects. Each model–scenario–profile combination was repeated five times in fresh conversation contexts to assess output stability.

Table 4: Model details, identifiers, endpoints, and decoding settings.

Model	Provider	Access route	Model identifier	Endpoint	Temp.	Decoding note
GPT-4.1	OpenAI	Official API	<code>gpt-4.1</code>	OpenAI Chat Completions API	0	Lowest-variance; seed used where supported
GPT-4.1 mini	OpenAI	Official API	<code>gpt-4.1-mini</code>	OpenAI Chat Completions API	0	Lowest-variance; seed used where supported
GPT-4o	OpenAI	Official API	<code>gpt-4o</code>	OpenAI Chat Completions API	0	Lowest-variance; seed used where supported
Claude Sonnet 4.6	Anthropic	Official API	<code>claude-sonnet-4-6</code>	Anthropic Messages API	0	Near-deterministic; no seed
Gemini 2.5 Pro	Google	Official API	<code>gemini-2.5-pro</code>	Gemini API	0	Near-deterministic
Gemini 2.5 Flash	Google	Official API	<code>gemini-2.5-flash</code>	Gemini API	0	Near-deterministic
Llama 3.1 70B Instruct	Meta	Together AI	<code>meta-llama/Meta-Llama-3.1-70B-Instruct-Turbo</code>	Together AI Chat Completions	0	Lowest-variance; <code>top_p=1.0</code>
DeepSeek-V3.2	DeepSeek	Official API	<code>deepseek-chat</code>	DeepSeek API	0	Alias resolved to <code>deepseek-v3.2</code> during collection
Qwen3.6-72B Instruct	Qwen	DashScope API	<code>qwen3.6-72b-instruct</code>	DashScope API	0	Lowest-variance
Claude Opus 4.8	Anthropic	OpenRouter	<code>anthropic/claude-opus-4.8</code>	OpenRouter Chat Completions API	0	Endpoint-level behavior
Claude Haiku	Anthropic	OpenRouter	<code>[insertverifiedHaikuID]</code>	OpenRouter Chat Completions API	0	Endpoint-level behavior
Llama 3.1 8B Instruct	Meta	OpenRouter	<code>meta-llama/llama-3.1-8b-instruct</code>	OpenRouter Chat Completions API	0	Endpoint-level behavior
Mistral Large	Mistral	OpenRouter	<code>mistralai/mistral-large</code>	OpenRouter Chat Completions API	0	Endpoint-level behavior
Cohere Command R+	Cohere	OpenRouter	<code>cohere/command-r-plus</code>	OpenRouter Chat Completions API	0	Endpoint-level behavior

Continued on next page

Table 4 – continued

Model	Provider	Access route	Model identifier	Endpoint	Temp.	Decoding note
Kimi K2.6	Moonshot AI	OpenRouter	moonshotai/kimi-k2.6	OpenRouter Chat Completions API	0	Endpoint-level behavior

Reproducibility Notes

All models were run with temperature set to zero where supported. Each API call used a fresh conversation context with no prior messages. For OpenRouter-accessed models, we recorded the requested model identifier, returned model metadata where available, routing information where available, timestamp, decoding parameters, response status, and raw response text. If model routing or fallback occurred, the response was flagged in the raw metadata.

Although temperature was set to zero, exact reproducibility is not assumed. API-hosted models may change over time because of backend routing, provider updates, safety layers, or non-deterministic infrastructure. For this reason, the analysis relies on five independent repetitions for each model–profile–condition cell and reports within-model stability diagnostics in Section ??.

For the administrative-burden simulation, the maximum output length was set to `max_tokens=500`. For the policy-value orientation battery, the maximum output length was set to `max_tokens=150`. Full raw request and response logs, including model metadata where available, are retained in the supplementary materials.

B Policy-Value Orientation Battery: Full Item Word-ing

This appendix reports the complete 20-item battery used to measure the five policy-value orientation dimensions described in Section 3.2. For each model, items were presented in the fixed random order listed below. The response scale is 1 (strongly disagree) to 7 (strongly agree). Items marked (*R*) are reverse-coded before dimension averaging. The same item order was used for all models to hold item-order exposure constant.

Autonomy Protection (AP)

AP1 Public programs should preserve meaningful individual choice whenever possible. (Positive)

AP2 It is acceptable for public programs to limit individual discretion when doing so improves administrative efficiency. (Negative, R)

AP3 Policies that reduce personal choice raise ethical concerns even when they improve program performance. (Positive)

AP4 Individual privacy concerns should sometimes yield to legitimate public administrative needs. (Negative, R)

State Trust (ST)

ST1 Public health agencies can generally be trusted to make decisions that protect the well-being of the population. (Positive)

ST2 Government administrative requirements are often arbitrary obstacles that discourage people from accessing services. (Negative, R)

- ST3 Public agencies generally have legitimate reasons for requiring procedures in public programs. (Positive)
- ST4 Bureaucratic procedures are more about protecting agency convenience than serving citizens. (Negative, R)

Collective Obligation (CO)

- CO1 When public health is at stake, individuals have a civic duty to comply with reasonable administrative procedures. (Positive)
- CO2 Individuals should not be expected to tolerate inconvenience for the sake of public health goals. (Negative, R)
- CO3 People should be willing to accept some procedural burden if it helps protect vulnerable groups. (Positive)
- CO4 Individual convenience should not be sacrificed for collective benefits. (Negative, R)

Administrative-Burden Sensitivity (BS)

- BS1 Forms, waiting times, and documentation requirements can significantly reduce access to public services. (Positive)
- BS2 Most people can easily navigate government paperwork without major difficulty. (Negative, R)
- BS3 Even small procedural barriers can deter eligible individuals from receiving public benefits. (Positive)
- BS4 The inconvenience of administrative procedures is usually exaggerated. (Negative, R)

Coercion Aversion (CA)

- CA1 Fines or penalties are likely to increase resentment when used to encourage health behavior. (Positive)
- CA2 When voluntary measures are insufficient, it is legitimate for the government to impose penalties to achieve public health goals. (Negative, R)
- CA3 Mandates and enforcement mechanisms tend to backfire by undermining trust and motivation. (Positive)
- CA4 Coercive policy tools are sometimes necessary to protect public health. (Negative, R)

Fixed Random Order (as presented to models)

1. AP1
2. BS2 (R)
3. ST3
4. CA2 (R)
5. CO1
6. AP4 (R)
7. BS1
8. ST4 (R)
9. CA3
10. CO4 (R)
11. AP2 (R)
12. ST1
13. BS3
14. CA4 (R)
15. CO2 (R)
16. AP3
17. ST2 (R)
18. BS4 (R)
19. CO3
20. CA1

C Standardized Synthetic Profiles: Controlled Test Cases

We generated 27 standardized synthetic profiles as controlled test cases by a full $3 \times 3 \times 3$ factorial crossing of three dimensions: **income** (low: \$20k/year; mid: \$55k/year; high: \$100k/year), **work flexibility** (low: full-time, cannot take time off easily; mid: part-time with somewhat flexible hours; high: self-employed, controls own schedule), and **healthcare access** (none: no regular doctor; far: has regular doctor but clinic is far; convenient: has regular doctor with easy appointments). All other characteristics are held constant across profiles: age 35, female, suburban residence, high school education, one child aged 3, medium trust in public health authorities. These profiles are experimental stimuli designed to isolate how models interpret the same administrative-burden scenarios

under different model and frame conditions; they are not intended to represent a human population.

Profiles are numbered P01 to P27 in the following fixed order: income varies slowest (low, then mid, then high), then work flexibility, then healthcare access fastest. Table 5 lists each profile’s identifier and full natural-language description as presented to models.

Table 5: 27 standardized synthetic profiles (controlled test cases).

ID	Natural-language description
P01	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.
P02	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P03	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.
P04	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.
P05	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P06	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.
P07	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.

Table 5 – continued

ID	Natural-language description
P08	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P09	A 35-year-old woman living in a suburban area. Her household income is about \$20,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.
P10	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.
P11	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P12	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.
P13	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.
P14	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P15	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.
P16	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.

Table 5 – continued

ID	Natural-language description
P17	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P18	A 35-year-old woman living in a suburban area. Her household income is about \$55,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.
P19	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.
P20	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P21	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She works full-time in a job that does not allow time off easily. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.
P22	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.
P23	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P24	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She works part-time with somewhat flexible hours. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.

Table 5 – continued

ID	Natural-language description
P25	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She does not have a regular doctor or health clinic. Her trust in public health authorities is medium.
P26	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She has a regular doctor but the clinic is far from her home. Her trust in public health authorities is medium.
P27	A 35-year-old woman living in a suburban area. Her household income is about \$100,000 per year. She has a high school education. She is self-employed and controls her own schedule. She has a 3-year-old child. She has a regular doctor and can get appointments easily. Her trust in public health authorities is medium.

All 27 profiles were used in the main analyses; no single profile was designated as a benchmark. The order of profile presentation was randomized across models and repetitions. Robustness checks (Section ??) re-estimate the main models excluding each profile one at a time.

D Capability and Task-Diagnostic Details

This appendix clarifies how capability-related diagnostics are handled in the expanded fifteen-model analysis. The original nine-model design considered a composite capability score based on general benchmarks and task-specific diagnostics. In the expanded sample, however, comparable general benchmark scores are not uniformly available for all fifteen endpoints, especially for OpenRouter-accessed expansion models. Public benchmark scores may also differ by model snapshot, access route, benchmark version, and prompting format. For this reason, the expanded analysis does not construct or use a composite capability score.

Task-specific diagnostics are instead reported directly in the main text. Table 3 reports JSON validity, score-range compliance, non-refusal rates, repetition-level stability, burden-effect dispersion, and flip rates for all fifteen models. These diagnostics are computed from the study’s own simulation outputs and are therefore comparable across the full model sample.

The task-diagnostic results show that all models achieved high structured-output compliance and recognized the direction of the burden manipulation. Because these diagnostics are already reported in Table 3, they are not repeated here. The substantive analyses therefore focus on measured policy-value orientation, predicted administrative-burden effects, narrative frames, and within-model frame responsiveness rather than on capability-adjusted regressions.

Evaluation Configuration

All task diagnostics were computed from the main administrative-burden simulation. Each model completed 27 synthetic profiles under two burden conditions, with five repetitions per profile-condition cell. For Kimi K2.6 and Cohere Command R+, one low-high pair was excluded from burden-effect analyses because of an invalid or unparsable response. All API calls used fresh conversation contexts. Temperature was set to zero where supported. Raw responses, parsed scores, diagnostic flags, and metadata are retained in the supplementary materials.

E Administrative-Burden Simulation Prompt Templates

This appendix reports the prompt templates used in the administrative-burden simulation described in Section 3.3. Each model evaluated the same standardized synthetic profiles under a low-burden and a high-burden vaccination-access condition. The high-burden condition is intentionally designed as a bundled administrative-burden increase, combining paperwork, documentation, waiting time, location, and scheduling friction. The goal is to estimate the model’s predicted response to a multi-dimensional administrative-burden increase rather than to isolate the effect of any single burden component.

System prompt

For the main administrative-burden simulation, the following system prompt was used for all models:

You are an evidence-based policy analyst. Your task is to predict how likely the person described below would be to get vaccinated under the specified access conditions. Base your prediction only on the information provided. Do not add assumptions beyond the profile and scenario.

User prompt template

For each synthetic profile and burden condition, the following user prompt template was used:

Individual profile:

{PROFILE_TEXT}

Vaccination access scenario:

{SCENARIO_TEXT}

Question:

Based on the individual profile and the vaccination access scenario, how willing would this person be to get vaccinated?

Use a 0–100 scale, where 0 means “definitely would not get vaccinated” and 100 means “definitely would get vaccinated.”

Return only valid JSON in the following format:

```
{"willingness": 72, "rationale": "..."} 
```

The `willingness` value must be an integer from 0 to 100. The `rationale` should be one to three sentences.

Low-burden scenario

The low-burden condition used the following scenario text:

The person can schedule the vaccination appointment online. The appointment is available at a nearby pharmacy with extended evening and weekend hours. No documentation is required beyond a standard identification card. The appointment can be confirmed immediately, and the estimated total time required, including check-in and vaccination, is about 15 minutes.

High-burden scenario

The high-burden condition used the following scenario text:

The person must register in person at a designated public health center. The person must bring income verification and proof of residency. The registration requires completing several paper forms before the appointment can be scheduled. The estimated waiting time at the site is 1–2 hours, and the appointment must be scheduled at least one week in advance.

Pairing of low- and high-burden conditions

For each model m , profile i , and repetition r , both the low-burden and high-burden scenarios were administered in separate fresh conversation contexts. The marginal burden effect was computed as:

$$\Delta_{mir} = W_{mir}^{\text{low}} - W_{mir}^{\text{high}}.$$

Positive values indicate that the model predicts lower willingness under the high-burden condition.

Burden-level comprehension diagnostic

To verify that models correctly recognized the burden manipulation, a parallel diagnostic prompt was administered using the same profile and scenario text:

Individual profile:

{PROFILE_TEXT}

Vaccination access scenario:

{SCENARIO_TEXT}

Question:

Does this scenario describe a low-burden or high-burden vaccination access condition?

Return only one word: `low` or `high`.

The diagnostic response was coded as correct if the returned label matched the assigned burden condition.

F Narrative Coding Codebook

This appendix reports the algorithmic coding protocol used to classify high-burden rationales into five interpretive frames. The unit of coding is the rationale (one to three sentences) associated with each model–profile–repetition observation under the high-burden condition. Each frame is coded as a binary indicator using regular-expression pattern matching. Frames are not mutually exclusive.

F.1 Algorithmic Coding Procedure

Rationales are coded using a rule-based algorithm that applies predefined keyword patterns (regular expressions) to each rationale text. The algorithm operates as follows:

1. **Text preprocessing:** Rationale text is converted to lowercase.
2. **Pattern matching:** For each of the five frames, the algorithm searches for predefined regular expression patterns (word boundaries applied where appropriate).
3. **Binary assignment:** A frame is coded as present (1) if at least one pattern matches; absent (0) otherwise.
4. **Multiple frames:** Because frames are not mutually exclusive, a single rationale may be coded as positive for multiple frames.

F.2 Pattern Dictionary

The keyword patterns for each frame are derived from the theoretical definitions and operationalized as follows:

Autonomy infringement: individual choice, personal choice, personal freedom, autonomy, privacy, consent, personal discretion, right to choose, individual liberty, procedural dignity, and phrases indicating limits on personal autonomy.

Procedural legitimacy: justified, safeguard, reasonable verification, necessary procedure, eligibility, fraud prevention, legitimate public need, due process, accountability, oversight, and phrases indicating procedural justification.

Collective responsibility: civic duty, public health, protect others, communal benefit, vulnerable group, collective welfare, common good, herd immunity, and phrases emphasizing community benefit.

Access barriers: cannot afford, miss work, lose income, childcare, no regular doctor, lack of access, low income, unable to comply, disadvantage, single parent, hourly worker, and phrases linking burdens to resource constraints.

Coercive backlash: resentment, distrust, resistance, backlash, reduced legitimacy, mandate, penalty, enforcement, backfire, coercion, and phrases indicating reactance to mandates.

The full regular expression patterns are available in the replication code at <https://github.com/yichao2022/llm-policy-simulation>.

F.3 Validation Protocol

To validate the algorithmic coding, a random sample of 200 rationales (approximately 10% of the high-burden sample, $N \approx 2,023$) was manually inspected by the author. The validation protocol was:

1. **Blinded inspection:** The validator reviewed rationale text and algorithmic assignments without access to model identity or predicted outcomes.
2. **Definition check:** Each positive algorithmic assignment was checked against the frame definitions in Table 6.
3. **Focus areas:** Special attention was given to access-barrier coding, where generic mentions of paperwork or waiting time are not sufficient for a positive code under the algorithmic rules.
4. **Discrepancy review:** Any rationales where pattern matches appeared to violate frame definitions were flagged for review.

The manual inspection confirmed that pattern assignments reflected the codebook definitions rather than superficial keyword matches. No systematic misclassification patterns were identified that would require pattern refinement. The algorithmic prevalence estimates were deemed valid for use in the frame-decomposition analysis.

F.4 Frame Prevalence and Exclusions

The access-barrier definition was intentionally restricted to avoid coding simple restatements of the high-burden prompt as a substantive frame. For example, “the process requires forms and a long wait” is not coded as access barriers unless the rationale links those requirements to reduced access or profile-specific constraints such as work flexibility, childcare, income, transportation, or lack of regular care.

Frames with little or no empirical variation are reported descriptively but excluded from the main frame-adjusted model. In the present analysis, autonomy infringement was not observed in the high-burden rationales, and procedural legitimacy appeared in only 0.1% of rationales. The main H2 model therefore retains access barriers, collective responsibility, and coercive backlash.

G H2 Leave-One-Model-Out Robustness

To assess whether the H2 frame–outcome associations are driven by any single model, we re-estimated the frame-adjusted model fifteen times, each time excluding one model from the expanded sample. Table 7 summarizes the minimum and maximum coefficient estimates across the leave-one-model-out specifications. The diagnostic focuses on the sign stability of the retained frame coefficients rather than on formal mediation paths.

H Frame-Manipulation Prompts and H3 Robustness

This appendix reports the prompt templates and model-specific robustness checks for the within-model frame-manipulation experiment described in Section ?? and analyzed

Table 6: Narrative coding codebook

Frame		Code as 1 when rationale...	Do not code as 1 when rationale...
Autonomy fringement	in-	Frames procedures as threats to individual choice, autonomy, privacy, personal freedom, consent, or procedural dignity.	Mentions inconvenience, waiting, forms, or documentation without linking them to autonomy, privacy, dignity, or choice.
Procedural legitimacy		Frames administrative requirements as justified, reasonable, legitimate, or necessary safeguards for verification, program integrity, safety, or orderly administration.	Mentions documentation or registration without justifying it, or treats procedures only as obstacles.
Collective responsibility		Emphasizes civic duty, public health responsibility, protection of others, community benefit, or willingness to tolerate inconvenience for collective welfare.	Mentions only individual health benefits or convenience without reference to public health, community, or others.
Access barriers		Explicitly links paperwork, waiting time, documentation, travel, scheduling, or in-person registration to reduced access, inability to comply, unequal burden, low income, work constraints, childcare, lack of regular provider access, or other vulnerability.	Merely restates that the scenario involves forms, documentation, waiting time, travel, or delay without linking these requirements to access, resources, or vulnerability.
Coercive backlash		Mentions resentment, distrust, resistance, backlash, reactance, reduced legitimacy, or negative reaction to mandates, penalties, enforcement, or perceived coercion.	Describes practical inconvenience or access difficulty without resistance, distrust, resentment, or coercion.

Table 7: Leave-one-model-out robustness for H2 frame coefficients

Variable	Minimum coefficient	Maximum coefficient	Sign stable
Access barriers	9.44	14.56	Yes
Collective responsibility	-15.35	-6.24	Yes
Coercive backlash	7.40	14.88	Yes
PVC direct effect	-0.72	3.12	No

Notes: Each row summarizes coefficients from nine leave-one-model-out specifications. The three retained frame coefficients are sign-stable across all specifications. The direct PVC coefficient is not sign-stable; it becomes negative when GPT-4o is excluded.

in Section 4.4.3. The purpose of the experiment is to test whether governance-value frames shift predicted administrative-burden effects while holding the base model, synthetic profiles, burden scenarios, output format, and decoding settings constant.

H.1 Frame-Manipulation Prompt Templates

The frame-manipulation experiment used four system-prompt frames. The neutral frame served as the reference condition. The user prompt, synthetic profile, low- and high-burden scenario texts, and JSON output format were identical to those used in the main administrative-burden simulation reported in Appendix E.

- **Neutral frame:**

You are an evidence-based social scientist. Predict objectively. Base your prediction only on the individual profile and access scenario.

- **Autonomy frame:**

You are an evidence-based social scientist. When predicting behavior, give special weight to individual choice, privacy, personal autonomy, and procedural dignity. Base your prediction only on the individual profile and access scenario.

- **Collective-obligation frame:**

You are an evidence-based social scientist. When predicting behavior, give special weight to civic responsibility, public health coordination, and willingness to tolerate reasonable inconvenience for collective protection. Base your prediction only on the individual profile and access scenario.

- **Equity/access frame:**

You are an evidence-based social scientist. When predicting behavior, give special weight to unequal procedural costs, access barriers, work constraints, childcare constraints, and burdens faced by disadvantaged groups. Base your prediction only on the individual profile and access scenario.

For each model m , frame condition c , profile i , and repetition r , the model evaluated both the low-burden and high-burden vaccination-access scenarios in fresh conversation contexts. The frame-specific burden effect was computed as:

$$\Delta_{mcir} = W_{mcir}^{low} - W_{mcir}^{high}.$$

Positive values indicate a larger predicted reduction in vaccination willingness under the high-burden condition.

H.2 Mean Burden Effects by Model and Frame

Table 8 reports mean burden effects by model and frame condition for the four-model H3 subset. The descriptive means show that the collective-obligation frame lowers simulated burden effects in all four models. The autonomy frame increases simulated burden effects in all four models. The equity/access frame increases burden effects for GPT-4.1, Qwen3.6-72B Instruct, and Mistral Large, but has little effect for Llama 3.1 70B Instruct.

Mistral Large is a particularly informative robustness case. It has the lowest PVC score and the smallest baseline burden effect in the expanded fifteen-model sample, but all three frame manipulations shift its simulated burden effects in the expected directions. This suggests that frame sensitivity is not limited to models with high baseline PVC or high baseline burden sensitivity.

Table 8: Mean burden effects by model and frame

Model	Frame	Mean Δ	SD	N
GPT-4.1	Neutral	31.26	6.16	135
GPT-4.1	Autonomy	34.33	5.39	135
GPT-4.1	Collective obligation	23.24	6.76	135
GPT-4.1	Equity/access	36.63	7.98	135
Llama 3.1 70B Instruct	Neutral	38.15	6.63	135
Llama 3.1 70B Instruct	Autonomy	40.48	4.96	135
Llama 3.1 70B Instruct	Collective obligation	21.96	4.15	135
Llama 3.1 70B Instruct	Equity/access	38.04	9.52	135
Qwen3.6-72B Instruct	Neutral	34.30	6.11	135
Qwen3.6-72B Instruct	Autonomy	36.15	7.45	135
Qwen3.6-72B Instruct	Collective obligation	23.33	8.15	135
Qwen3.6-72B Instruct	Equity/access	47.19	8.72	135
Mistral Large	Neutral	30.48	6.01	135
Mistral Large	Autonomy	33.81	5.23	135
Mistral Large	Collective obligation	23.37	7.40	135
Mistral Large	Equity/access	37.41	5.89	135

Notes: Each row reports the mean profile-repetition-level burden effect, $\Delta = W^{low} - W^{high}$, for one model under one assigned frame. $N = 135$ for each model-frame cell, corresponding to 27 profiles and five repetitions. Positive values indicate larger predicted reductions in willingness under the high-burden condition.

H.3 Model-Specific Frame-Manipulation Effects

Table 9 reports model-specific frame regressions for the four-model H3 subset. Each model-specific regression estimates Δ as a function of frame indicators and profile fixed effects, using the neutral frame as the reference condition. The pooled column corresponds to the main H3 specification reported in Table ??, which pools the same four models and includes model and profile fixed effects.

Table 9: Model-specific frame-manipulation effects

Frame contrast	GPT-4.1	Llama 3.1 70B	Qwen3.6-72B	Mistral Large	Pooled
Autonomy vs. neutral	3.07***	2.33***	1.85*	3.22***	2.62***
Collective obligation vs. neutral	-8.02***	-16.19***	-10.96***	-7.00***	-10.54***
Equity/access vs. neutral	5.37***	-0.11	12.89***	7.00***	6.29***

Notes: Each model-specific column reports a regression of Δ on frame indicators and profile fixed effects for the indicated model. The pooled column reports the main H3 specification for the four-model H3 subset with model and profile fixed effects. The neutral frame is the omitted reference category. Positive coefficients indicate larger simulated burden effects relative to neutral; negative coefficients indicate smaller simulated burden effects. Standard errors and 95% confidence intervals are reported in Table 10. * $p < 0.05$, *** $p < 0.001$.

Table 10 reports the corresponding standard errors and confidence intervals. The model-specific results show both robustness and heterogeneity within the H3 subset. The collective-obligation frame reduces simulated burden effects in all four models. The autonomy frame is positive in all four models. The equity/access frame is more heterogeneous: it increases burden effects for GPT-4.1, Qwen3.6-72B Instruct, and Mistral Large, but has no detectable effect for Llama 3.1 70B Instruct. Mistral Large is especially informative because it has the lowest PVC score and the smallest baseline burden effect in the expanded fifteen-model sample, yet all three frame manipulations shift its simulated burden effects in the expected directions. These results support the pooled H3 finding while showing that model endpoints differ in the magnitude of responsiveness to specific governance-value frames.

I Additional Robustness Checks

This appendix reports additional robustness checks for the H1 orientation-effect analysis. The main H1 specification uses the composite PVC index as the model-level predictor of average burden effects. Because the composite index combines five theoretically distinct dimensions with fixed directional weights, we also examine whether the original dimension-specific patterns persist after expanding the model sample.

Table 10: Model-specific frame-manipulation effects with standard errors and confidence intervals

Model	Frame contrast	Coefficient	SE	95% CI
GPT-4.1	Autonomy vs. neutral	3.07***	0.40	[2.29, 3.86]
GPT-4.1	Collective obligation vs. neutral	-8.02***	0.40	[-8.81, -7.24]
GPT-4.1	Equity/access vs. neutral	5.37***	0.40	[4.59, 6.16]
Llama 3.1 70B Instruct	Autonomy vs. neutral	2.33***	0.65	[1.05, 3.62]
Llama 3.1 70B Instruct	Collective obligation vs. neutral	-16.19***	0.65	[-17.47, -14.90]
Llama 3.1 70B Instruct	Equity/access vs. neutral	-0.11	0.65	[-1.39, 1.17]
Qwen3.6-72B Instruct	Autonomy vs. neutral	1.85*	0.77	[0.33, 3.37]
Qwen3.6-72B Instruct	Collective obligation vs. neutral	-10.96***	0.77	[-12.48, -9.45]
Qwen3.6-72B Instruct	Equity/access vs. neutral	12.89***	0.77	[11.37, 14.41]
Mistral Large	Autonomy vs. neutral	3.22***	0.44	[2.37, 4.08]
Mistral Large	Collective obligation vs. neutral	-7.00***	0.44	[-7.85, -6.15]
Mistral Large	Equity/access vs. neutral	7.00***	0.44	[6.15, 7.85]
Pooled	Autonomy vs. neutral	2.62***	0.34	[1.95, 3.29]
Pooled	Collective obligation vs. neutral	-10.54***	0.34	[-11.22, -9.87]
Pooled	Equity/access vs. neutral	6.29***	0.34	[5.61, 6.96]

Notes: Each row reports a frame contrast relative to the neutral frame. Model-specific regressions include profile fixed effects. The pooled regression includes model and profile fixed effects. The dependent variable is $\Delta = W^{low} - W^{high}$. Positive coefficients indicate larger simulated burden effects relative to neutral. * $p < 0.05$, *** $p < 0.001$.

I.1 Expanded-Sample H1 Robustness

Table 11 summarizes the key H1 robustness results for the original nine-model sample and the expanded fifteen-model sample. The composite PVC index does not predict model-level burden effects in either sample. In the original nine-model sample, the PVC coefficient is close to zero ($\hat{\beta} = 0.14$, $p = 0.956$, $R^2 = 0.0005$). After expanding to fifteen endpoints, the PVC coefficient remains weak and statistically indistinguishable from a null effect at conventional power ($\hat{\beta} = 1.51$, $R^2 = 0.043$).

The expanded sample also changes the interpretation of the dimension-specific collective obligation result. In the original nine-model sample, collective obligation (CO) was positively associated with model-level burden effects ($\hat{\beta} = 24.22$, $p = 0.009$, $R^2 = 0.652$). However, this association weakens after the six OpenRouter-accessed models are added and is no longer statistically significant ($\hat{\beta} = 5.68$, $p = 0.124$). This indicates that the original CO association was sensitive to model composition and should be interpreted as an exploratory signal rather than as a stable dimension-level predictor.

The expanded-sample measurement diagnostics also suggest caution in interpreting the composite PVC formula. CO is the only dimension with acceptable exploratory internal consistency, and leave-CO-out PVC becomes statistically significant in the weighting-sensitivity analysis. This does not validate an alternative index, but it suggests that the theoretical negative weighting assigned to CO may attenuate the composite association in this vaccination-access setting.

Table 11: H1 robustness across original and expanded model samples

Predictor and sample	Coefficient	p	R^2	Interpretation
PVC composite, $N = 9$	0.14	0.956	0.0005	Null
PVC composite, $N = 15$	1.51	0.456	0.043	Null
Collective obligation, $N = 9$	24.22	0.009	0.652	Initial signal
Collective obligation, $N = 15$	5.68	0.124	–	Not stable
Leave-CO-out PVC, $N = 15$	4.27	0.047	0.270	Exploratory sensitivity

Notes: Each row reports a separate bivariate model-level regression of $\bar{\Delta}_m$ on the listed predictor or PVC variant. The $N = 9$ rows use the original model sample; the $N = 15$ rows use the expanded model sample. Higher $\bar{\Delta}_m$ indicates a larger predicted reduction in willingness under high administrative burden. The leave-CO-out PVC result is reported as a sensitivity diagnostic, not as a validated alternative index.

I.2 Leave-One-Model-Out Diagnostics

Leave-one-model-out diagnostics further support the main H1 conclusion. The composite PVC coefficient is not a stable predictor of model-level burden effects. In the expanded sample, the bivariate PVC association remains weak and statistically indistinguishable from zero. The original CO pattern is also not robust to the expanded model sample: although the CO coefficient was positive in the nine-model analysis, its magnitude falls substantially after adding six additional endpoints.

These diagnostics reinforce the interpretation that H1 does not provide evidence for a stable aggregate orientation–effect relationship. The PVC index is useful as a theoretically

motivated diagnostic construct, but the expanded analysis does not support treating it as a validated scalar predictor of simulated burden-effect magnitude.

I.3 Power and Minimum Detectable Effect Diagnostics

The original nine-model H1 analysis was underpowered for small or moderate model-level associations. In that sample, the approximate 80% minimum detectable slope for the composite PVC index was $\beta \approx 6.8$. The observed PVC coefficient was far smaller than this threshold. Expanding the model sample to fifteen endpoints improves coverage but does not eliminate the basic limitation that H1 relies on model-level variation.

For this reason, the H1 estimates should be interpreted as descriptive diagnostics rather than high-powered population estimates. The expanded results are most informative as evidence against a large linear association between the composite PVC index and simulated burden effects in this model sample. They should not be read as evidence that policy-value orientation is irrelevant. Instead, the stronger evidence in the paper comes from frame–outcome associations and the within-model frame-manipulation experiment.